

ipython\_notebooks (/github/pybokeh/ipython\_notebooks/tree/master) / pandas (/github/pybokeh/ipython\_notebooks/tree/master/pandas)

## Python Pandas (<http://pandas.pydata.org/pandas-docs/stable/>) Cheat Sheet

As a data analyst, these are common tasks I perform using Pandas

- [Pandas display output options](#)
- [Quick summary statistics, # of rows, columns, etc](#)
- [Sorting More Than One Column](#)
- [Removing duplicates](#)
- [Creating a column based on values from another column](#)
- [Removing a column](#)
- [Replacing values in a series](#)
- [Reading CSV or Excel files \(<http://pandas.pydata.org/pandas-docs/stable/io.html>\)](#)
- [Renaming axis](#)
- [Renaming indexes or columns](#)
- [Binning data into intervals](#)
- [Grouping data](#)
- [Sub-selecting or Slicing a Data Frame](#)
- [Boolean indexing](#)
- [Obtaining columns with partial column labels](#)
- [Getting value counts](#)
- [Getting cumulative sum](#)
- [Pivot table](#)
- [Creating a new column based on a grouping using transform\(\) method](#)
- [Pivot table Percent of Rows](#)
- [Pivot table Percent of Columns](#)
- [Transpose a data frame](#)
- [Converting index to a column](#)
- [Converting column to index](#)
- [How to add or fill in missing dates](#)
- [How to connect to an ODBC data source \(Windows\)](#)
- [How to convert data in wide format to long format using melt\(\)](#)
- [How to convert data in long format data to wide format using pivot\(\)](#)
- [Using category data type to control sort order](#)
- [Merging 2 data frames using merge\(\)](#)
- [Finding rows containing data with missing values](#)
- [Converting a data type of a column in a data frame](#)
- [Plotting data frames using MATPLOTLIB ver 1.5+](#)
- [Plotting data frame directly and creating sub-plots](#)
- [Creating a Bokeh chart from a Data Frame](#)
- [Making plotly charts directly from data frame](#)
- [Method chaining](#)
- [Sending Pandas data frame to R using rpy2 IPython notebook extension](#)
- [Python clone of R's dplyr \(\[http://nbviewer.jupyter.org/github/pybokeh/jupyter\\\_notebooks/blob/master/dplython/dplython\\\_example.ipynb\]\(http://nbviewer.jupyter.org/github/pybokeh/jupyter\_notebooks/blob/master/dplython/dplython\_example.ipynb\)\)](#)
- [BONUS #1: A HUGE list of python and pandas snippets by Chris Albon \(<http://chrisalbon.com/>\)](#)
- [BONUS #2: More goodies from a major pandas contributor, Tom Augspurger \(<http://tomaugspurger.github.io/>\)](#)

Setting display output options ([http://pandas.pydata.org/pandas-docs/stable/generated/pandas.set\\_option.html](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.set_option.html))

[\[back to top\]](#)

In [70]:

```
import pandas as pd
pd.set_option("display.max_rows",1000)    # or pd.options.display.max_rows=1000
pd.set_option("display.max_columns",20)   # or pd.options.display.max_columns=20
pd.set_option('precision',7)
pd.set_option('large_repr', 'truncate')
```

**Quick summary statistics using `df.describe()` and data types using `df.info()`. Also check out [pandas-summary](https://github.com/mouradmourafiq/pandas-summary) (<https://github.com/mouradmourafiq/pandas-summary>) - an extension for data frame's describe() method.**

[\[back to top\]](#)

In [4]:

```
import pandas as pd
data = pd.DataFrame({'group': ['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'c'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[4]:

|   | group | ounces |
|---|-------|--------|
| 0 | a     | 4.0    |
| 1 | a     | 3.0    |
| 2 | a     | 12.0   |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |

In [5]:

```
data.describe()
```

Out[5]:

|       | ounces    |
|-------|-----------|
| count | 9.000000  |
| mean  | 6.055556  |
| std   | 2.855307  |
| min   | 3.000000  |
| 25%   | 4.000000  |
| 50%   | 6.000000  |
| 75%   | 7.500000  |
| max   | 12.000000 |

**df.info() shows data types, number of rows and columns, and memory usage of your data frame**

In [6]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9 entries, 0 to 8
Data columns (total 2 columns):
group      9 non-null object
ounces     9 non-null float64
dtypes: float64(1), object(1)
memory usage: 216.0+ bytes
```

## Sorting More Than One Column

[\[back to top\]](#)

**Just pass a list of columns you want to sort and pass corresponding list of True/False to the ascending parameter**

In [16]:

```
data.sort_values(by=['group', 'ounces'], ascending=[False, True], inplace=True)
data
```

Out[16]:

|   | group | ounces |
|---|-------|--------|
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 1 | a     | 3.0    |
| 0 | a     | 4.0    |
| 2 | a     | 12.0   |

## Removing duplicates

[\[back to top\]](#)

In [19]:

```
import pandas as pd
data = pd.DataFrame({'k1': ['one'] * 3 + ['two'] * 4, 'k2': [3, 2, 1, 3, 3, 4, 4]})
```

In [20]:

```
data.sort_values(by='k2')
data
```

Out[20]:

|   | k1  | k2 |
|---|-----|----|
| 0 | one | 3  |
| 1 | one | 2  |
| 2 | one | 1  |
| 3 | two | 3  |
| 4 | two | 3  |
| 5 | two | 4  |
| 6 | two | 4  |

In [21]:

```
data.drop_duplicates() # by default, duplicate is defined by all columns
```

Out[21]:

|   | k1  | k2 |
|---|-----|----|
| 0 | one | 3  |
| 1 | one | 2  |
| 2 | one | 1  |
| 3 | two | 3  |
| 5 | two | 4  |

## Define duplicates by column name(s):

In [22]:

```
data.drop_duplicates(subset='k1') # duplicate in column k1 only
```

Out[22]:

|   | k1  | k2 |
|---|-----|----|
| 0 | one | 3  |
| 3 | two | 3  |

## Creating a new column based on values from another column

[\[back to top\]](#)

In [5]:

```
import pandas as pd
data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon', 'Pastrami','corned beef', 'Bacon', 'pastrami', 'honey ham','nova lox'],
                    'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[5]:

|   | food        | ounces |
|---|-------------|--------|
| 0 | bacon       | 4.0    |
| 1 | pulled pork | 3.0    |
| 2 | bacon       | 12.0   |
| 3 | Pastrami    | 6.0    |
| 4 | corned beef | 7.5    |
| 5 | Bacon       | 8.0    |
| 6 | pastrami    | 3.0    |
| 7 | honey ham   | 5.0    |
| 8 | nova lox    | 6.0    |

Suppose you wanted to add a column indicating the type of animal that each food came from. Let's write down a mapping of each distinct meat type to the kind of animal using a dictionary and a function:

In [7]:

```
meat_to_animal = {
    'bacon': 'pig',
    'pulled pork': 'pig',
    'pastrami': 'cow',
    'corned beef': 'cow',
    'honey ham': 'pig',
    'nova lox': 'salmon'
}

def meat2animal(series):
    if series["food"]=="bacon":
        return 'pig'
    elif series["food"]=="pulled pork":
        return 'pig'
    elif series["food"]=="pastrami":
        return 'cow'
    elif series["food"]=="corned beef":
        return 'cow'
    elif series["food"]=="honey ham":
        return 'pig'
    else:
        return 'salmon'
```

In [8]:

```
data['animal'] = data['food'].map(str.lower).map(meat_to_animal)
data
```

Out[8]:

|   | food        | ounces | animal |
|---|-------------|--------|--------|
| 0 | bacon       | 4.0    | pig    |
| 1 | pulled pork | 3.0    | pig    |
| 2 | bacon       | 12.0   | pig    |
| 3 | Pastrami    | 6.0    | cow    |
| 4 | corned beef | 7.5    | cow    |
| 5 | Bacon       | 8.0    | pig    |
| 6 | pastrami    | 3.0    | cow    |
| 7 | honey ham   | 5.0    | pig    |
| 8 | nova lox    | 6.0    | salmon |

Or we could use `apply()` and apply the `meat2animal()` function to the new column:

In [9]:

```
# axis=1 means to apply the function for each row, but I prefer to use axis='columns' instead
data['animal2'] = data.apply(meat2animal,axis='columns')
data
```

Out[9]:

|   | food        | ounces | animal | animal2 |
|---|-------------|--------|--------|---------|
| 0 | bacon       | 4.0    | pig    | pig     |
| 1 | pulled pork | 3.0    | pig    | pig     |
| 2 | bacon       | 12.0   | pig    | pig     |
| 3 | Pastrami    | 6.0    | cow    | salmon  |
| 4 | corned beef | 7.5    | cow    | cow     |
| 5 | Bacon       | 8.0    | pig    | salmon  |
| 6 | pastrami    | 3.0    | cow    | cow     |
| 7 | honey ham   | 5.0    | pig    | pig     |
| 8 | nova lox    | 6.0    | salmon | salmon  |

We can also use `dataframe.assign()` (<http://pandas.pydata.org/pandas-docs/stable/dsintro.html#dsintro-chained-assignment>) function which was added in pandas version 0.16

In [4]:

```
import pandas as pd
import numpy as np

df = pd.DataFrame({'data1' : np.random.randn(5),
                  'data2' : np.random.randn(5)})

df.assign(ratio = df['data1'] / df['data2'])
```

Out[4]:

|   | data1     | data2     | ratio     |
|---|-----------|-----------|-----------|
| 0 | 0.246394  | -0.496658 | -0.496105 |
| 1 | 1.690199  | 1.245658  | 1.356873  |
| 2 | -0.132093 | 0.834331  | -0.158322 |
| 3 | -0.046662 | 0.717323  | -0.065050 |
| 4 | -0.355658 | 0.439520  | -0.809196 |

[\[back to top\]](#)

## Removing or dropping a column

[\[back top top\]](#)

In [10]:

```
data.drop('animal2', axis='columns', inplace=True) # dropping/removing a column
data
```

Out[10]:

|   | food        | ounces | animal |
|---|-------------|--------|--------|
| 0 | bacon       | 4.0    | pig    |
| 1 | pulled pork | 3.0    | pig    |
| 2 | bacon       | 12.0   | pig    |
| 3 | Pastrami    | 6.0    | cow    |
| 4 | corned beef | 7.5    | cow    |
| 5 | Bacon       | 8.0    | pig    |
| 6 | pastrami    | 3.0    | cow    |
| 7 | honey ham   | 5.0    | pig    |
| 8 | nova lox    | 6.0    | salmon |

## Replacing Values in a Series

[\[back to top\]](#)

In [30]:

```
import pandas as pd
data = pd.Series([1., -999., 2., -999., -1000., 3.])
data
```

Out[30]:

```
0      1
1   -999
2       2
3   -999
4  -1000
5       3
dtype: float64
```

**If you want to replace -999 with NaN:**

In [32]:

```
data.replace(-999, np.nan, inplace=True)
data
```

Out[32]:

```
0      1
1     NaN
2       2
3     NaN
4  -1000
5       3
dtype: float64
```

**If you want to replace multiple values at once, you instead pass a list then the substitute value:**

In [35]:

```
import pandas as pd
data = pd.Series([1., -999., 2., -999., -1000., 3.])
data
```

Out[35]:

```
0      1
1   -999
2       2
3   -999
4  -1000
5       3
dtype: float64
```

In [37]:

```
data.replace([-999, -1000], np.nan, inplace=True)
data
```

Out[37]:

```
0      1
1     NaN
2       2
3     NaN
4     NaN
5       3
dtype: float64
```

## Renaming Axis Indexes

[\[back to top\]](#)

In [47]:

```
import pandas as pd
data = pd.DataFrame(np.arange(12).reshape((3, 4)),index=['Ohio', 'Colorado', 'New York'],columns=['one', 'two', 'three', 'four'])
data
```

Out[47]:

|          | one | two | three | four |
|----------|-----|-----|-------|------|
| Ohio     | 0   | 1   | 2     | 3    |
| Colorado | 4   | 5   | 6     | 7    |
| New York | 8   | 9   | 10    | 11   |

3 rows × 4 columns

In [48]:

```
data.index.map(str.upper)
```

Out[48]:

```
array(['OHIO', 'COLORADO', 'NEW YORK'], dtype=object)
```

In [51]:

```
data.index = data.index.map(str.upper)
data
```

Out[51]:

|          | one | two | three | four |
|----------|-----|-----|-------|------|
| OHIO     | 0   | 1   | 2     | 3    |
| COLORADO | 4   | 5   | 6     | 7    |
| NEW YORK | 8   | 9   | 10    | 11   |

3 rows × 4 columns

If you want to create a transformed version of a data set without modifying the original, a useful method is `rename`:

In [65]:

```
data.rename(index=str.title, columns=str.upper, inplace=True) # str.title means to make the 1st letter capitalized only
data
```

Out[65]:

|          | ONE | TWO | THREE | FOUR |
|----------|-----|-----|-------|------|
| Ohio     | 0   | 1   | 2     | 3    |
| Colorado | 4   | 5   | 6     | 7    |
| New York | 8   | 9   | 10    | 11   |

3 rows × 4 columns

## Renaming Indexes or Columns

[\[back to top\]](#)

`rename()` can be used in conjunction with a dict-like object providing new values for a subset of the axis labels:

In [46]:

```
import pandas as pd
data = pd.DataFrame(np.arange(12).reshape((3, 4)),index=['Ohio', 'Colorado', 'New York'],columns=['one', 'two', 'three', 'four'])
data
```

Out[46]:

|          | one | two | three | four |
|----------|-----|-----|-------|------|
| Ohio     | 0   | 1   | 2     | 3    |
| Colorado | 4   | 5   | 6     | 7    |
| New York | 8   | 9   | 10    | 11   |

In [47]:

```
data.rename(index={'Ohio': 'INDIANA'},columns={'three': 'peekaboo'},inplace=True)
data
```

Out[47]:

|          | one | two | peekaboo | four |
|----------|-----|-----|----------|------|
| INDIANA  | 0   | 1   | 2        | 3    |
| Colorado | 4   | 5   | 6        | 7    |
| New York | 8   | 9   | 10       | 11   |

You can also apply str functions to modify the index or column labels

In [48]:

```
data.rename(index=str.title, columns=str.upper, inplace=True) # str.title means to make the 1st letter capitalized only
data
```

Out[48]:

|          | ONE | TWO | PEEKABOO | FOUR |
|----------|-----|-----|----------|------|
| Indiana  | 0   | 1   | 2        | 3    |
| Colorado | 4   | 5   | 6        | 7    |
| New York | 8   | 9   | 10       | 11   |

## Binning Data Into Intervals

[\[back to top\]](#)

In [71]:

```
ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

Let's divide these into bins of 18 to 25, 26 to 35, 35 to 60, and finally 60 and older. To do so, you have to use `cut()` function:

In [73]:

```
bins = [18, 25, 35, 60, 100]
```

In [75]:

```
cats = pd.cut(ages, bins)
cats
```

Out[75]:

```
(18, 25]
(18, 25]
(18, 25]
(25, 35]
(18, 25]
(18, 25]
(35, 60]
(25, 35]
(60, 100]
(35, 60]
(35, 60]
(25, 35]
Levels (4): Index(['(18, 25]', '(25, 35]', '(35, 60]', '(60, 100]'], dtype=object)
```

In [77]:

```
cats.labels
```

Out[77]:

```
array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int32)
```

In [78]:

```
cats.levels
```

Out[78]:

```
Index(['(18, 25]', '(25, 35]', '(35, 60]', '(60, 100]'], dtype='object')
```



In [79]:

```
pd.value_counts(cats)
```

Out[79]:

```
(18, 25]    5
(35, 60]    3
(25, 35]    3
(60, 100]   1
dtype: int64
```

If you want the right value to be exclusive in the intervals:

In [80]:

```
pd.cut(ages, [18, 26, 36, 61, 100], right=False)
```

Out[80]:

```
[18, 26)
[18, 26)
[18, 26)
[26, 36)
[18, 26)
[18, 26)
[36, 61)
[26, 36)
[61, 100)
[36, 61)
[36, 61)
[26, 36)
Levels (4): Index(['[18, 26)', '[26, 36)', '[36, 61)', '[61, 100)'], dtype=object)
```

You can also pass your own bin names by passing a list or array to the labels option:

In [83]:

```
group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'Senior']
pd.cut(ages, bins, labels=group_names)
pd.value_counts(pd.cut(ages, bins, labels=group_names))
```

Out[83]:

```
Youth      5
YoungAdult  3
MiddleAged  3
Senior     1
dtype: int64
```

[\[back to top\]](#)

## Grouping Data ( see also [value count\(\)](#) and [pivot table\(\)](#) )

[\[back to top\]](#)

In [1]:

```
import pandas as pd

df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                   'key2' : ['one', 'two', 'one', 'two', 'one'],
                   'data1' : np.random.randn(5),
                   'data2' : np.random.randn(5)})

df
```

Out[1]:

|   | data1     | data2    | key1 | key2 |
|---|-----------|----------|------|------|
| 0 | -1.206295 | 0.277025 | a    | one  |
| 1 | -0.210670 | 1.700554 | a    | two  |
| 2 | 1.090656  | 0.141390 | b    | one  |
| 3 | -0.755093 | 1.402054 | b    | two  |
| 4 | 0.251278  | 0.767201 | a    | one  |

5 rows × 4 columns

In [3]:

```
grouped = df['data1'].groupby(df['key1'])
grouped.mean()
```

Out[3]:

```
key1
a      -0.388562
b       0.167782
dtype: float64
```

## Sub-selecting or Slicing a Data Frame

[\[back to top\]](#)
Filtering by label name: [\[loc\]](#)Filtering by index row and/or column: [\[iloc\]](#)

In [1]:

```
import numpy as np
import pandas as pd

dates = pd.date_range('20130101', periods=6)
df = pd.DataFrame(np.random.randn(6,4), index=dates, columns=list('ABCD'))
df
```

Out[1]:

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-01 | 0.618255  | 0.699026  | -0.857985 | -0.167018 |
| 2013-01-02 | 1.193699  | 0.304388  | 0.036477  | 0.073046  |
| 2013-01-03 | -1.409751 | 0.599982  | 1.359995  | 0.901001  |
| 2013-01-04 | 0.165882  | -0.156601 | -0.764553 | 0.665824  |
| 2013-01-05 | 0.190956  | -0.515755 | -0.866012 | -0.908422 |
| 2013-01-06 | 0.849617  | -1.473412 | 0.013182  | 0.443061  |

## Getting first n rows of data frame using index slicing syntax

In [2]:

```
df[0:3] # get first 3 rows of the data frame
```

Out[2]:

|            | A        | B         | C        | D         |
|------------|----------|-----------|----------|-----------|
| 2013-01-01 | 1.495158 | 2.136794  | 0.034707 | -0.128909 |
| 2013-01-02 | 0.562039 | -1.150613 | 0.265513 | 0.093981  |
| 2013-01-03 | 0.973170 | 0.877070  | 0.682884 | -0.026271 |

## Slicing based on data frame's index range

In [3]:

```
df['20130102':'20130104'] # get rows by index range
```

Out[3]:

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-02 | 0.562039  | -1.150613 | 0.265513  | 0.093981  |
| 2013-01-03 | 0.973170  | 0.877070  | 0.682884  | -0.026271 |
| 2013-01-04 | -1.146283 | 1.350325  | -0.040049 | 1.069154  |

## Slicing based on column labels/names using loc

[\[view df\]](#)

In [4]:

```
df.loc[:,['A','B']] # syntax is: df.loc[rows_index, cols_index]
```

Out[4]:

|            | A         | B         |
|------------|-----------|-----------|
| 2013-01-01 | 1.495158  | 2.136794  |
| 2013-01-02 | 0.562039  | -1.150613 |
| 2013-01-03 | 0.973170  | 0.877070  |
| 2013-01-04 | -1.146283 | 1.350325  |
| 2013-01-05 | 0.919800  | -0.058590 |
| 2013-01-06 | 1.286808  | -0.357197 |

### Slicing based on row index label and column label combined using loc

In [5]:

```
df.loc['20130102':'20130104',['A','B']]
```

Out[5]:

|            | A         | B         |
|------------|-----------|-----------|
| 2013-01-02 | 0.562039  | -1.150613 |
| 2013-01-03 | 0.973170  | 0.877070  |
| 2013-01-04 | -1.146283 | 1.350325  |

### Slicing based on index position of the row or column using iloc

[\[view df\]](#)

In [6]:

```
df.iloc[3] # returns 4th row (index=3) of the data frame
```

Out[6]:

```
A    -1.146283
B     1.350325
C    -0.040049
D     1.069154
Name: 2013-01-04 00:00:00, dtype: float64
```

In [7]:

```
df.iloc[3:5,0:2] # returns specific range of rows and columns of the data frame
```

Out[7]:

|            | A         | B         |
|------------|-----------|-----------|
| 2013-01-04 | -1.146283 | 1.350325  |
| 2013-01-05 | 0.919800  | -0.058590 |

In [9]:

```
df.iloc[[1,5],[0,2]] # returns specific rows and columns using lists containing columns or row indexes
```

Out[9]:

|            | A        | C         |
|------------|----------|-----------|
| 2013-01-02 | 0.562039 | 0.265513  |
| 2013-01-06 | 1.286808 | -0.205763 |

In [10]:

```
df.iloc[1:3,:] # returning specific rows and returning all columns
```

Out[10]:

|            | A        | B         | C        | D         |
|------------|----------|-----------|----------|-----------|
| 2013-01-02 | 0.562039 | -1.150613 | 0.265513 | 0.093981  |
| 2013-01-03 | 0.973170 | 0.877070  | 0.682884 | -0.026271 |

In [11]:

```
df.iloc[:,1:3] # returning all rows and specific columns
```

Out[11]:

|            | B         | C         |
|------------|-----------|-----------|
| 2013-01-01 | 2.136794  | 0.034707  |
| 2013-01-02 | -1.150613 | 0.265513  |
| 2013-01-03 | 0.877070  | 0.682884  |
| 2013-01-04 | 1.350325  | -0.040049 |
| 2013-01-05 | -0.058590 | 0.083294  |
| 2013-01-06 | -0.357197 | -0.205763 |

In [12]:

```
df.iloc[1,1] # getting specific scalar/single value
```

Out[12]:

```
-1.1506133975623973
```

## **Boolean Indexing (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#boolean-indexing>)**

[\[view df\]](#)

The boolean operators are: | for or, & for and, and ~ for not. These **must** be grouped by using parentheses.

In [33]:

```
df[df.A > -0.5] # or df[df["A"] > -0.5], this syntax works when there is a space in the column name
```

Out[33]:

|            | A         | B         | C         | D        |
|------------|-----------|-----------|-----------|----------|
| 2013-01-03 | 2.351035  | 0.085318  | -2.265766 | 0.432268 |
| 2013-01-04 | 0.447390  | 0.205224  | -0.614702 | 0.212125 |
| 2013-01-05 | -0.455682 | -2.765288 | 0.475253  | 0.432910 |
| 2013-01-06 | 1.000494  | -0.462591 | -0.638053 | 0.481087 |

or I've seen some people do boolean filtering by passing "criteria" variable to the data frame like so:

In [34]:

```
criteria = df.A > -0.5
df[criteria]
```

Out[34]:

|            | A         | B         | C         | D        |
|------------|-----------|-----------|-----------|----------|
| 2013-01-03 | 2.351035  | 0.085318  | -2.265766 | 0.432268 |
| 2013-01-04 | 0.447390  | 0.205224  | -0.614702 | 0.212125 |
| 2013-01-05 | -0.455682 | -2.765288 | 0.475253  | 0.432910 |
| 2013-01-06 | 1.000494  | -0.462591 | -0.638053 | 0.481087 |

In [2]:

```
df2 = df.copy()
df2['E']=['one', 'one', 'two', 'three', 'four', 'three']
df2
```

Out[2]:

|            | A         | B         | C         | D         | E     |
|------------|-----------|-----------|-----------|-----------|-------|
| 2013-01-01 | 0.618255  | 0.699026  | -0.857985 | -0.167018 | one   |
| 2013-01-02 | 1.193699  | 0.304388  | 0.036477  | 0.073046  | one   |
| 2013-01-03 | -1.409751 | 0.599982  | 1.359995  | 0.901001  | two   |
| 2013-01-04 | 0.165882  | -0.156601 | -0.764553 | 0.665824  | three |
| 2013-01-05 | 0.190956  | -0.515755 | -0.866012 | -0.908422 | four  |
| 2013-01-06 | 0.849617  | -1.473412 | 0.013182  | 0.443061  | three |

In [36]:

```
df2[df2['E'].isin(['two','four'])] # read as "return rows where column E contains two or four"
```

Out[36]:

|            | A         | B         | C         | D        | E    |
|------------|-----------|-----------|-----------|----------|------|
| 2013-01-03 | 2.351035  | 0.085318  | -2.265766 | 0.432268 | two  |
| 2013-01-05 | -0.455682 | -2.765288 | 0.475253  | 0.432910 | four |

We can also do partial string matching. So let's say you don't know the exact spelling a word is you want to match, you can do this:

In [3]:

```
df2[df2.E.str.contains("tw|ou")]
```

Out[3]:

|            | A         | B         | C         | D         | E    |
|------------|-----------|-----------|-----------|-----------|------|
| 2013-01-03 | -1.409751 | 0.599982  | 1.359995  | 0.901001  | two  |
| 2013-01-05 | 0.190956  | -0.515755 | -0.866012 | -0.908422 | four |

Using ~ to do a "NOT"

In [38]:

```
df2[~df2['E'].isin(['two','four'])] # column E containing values not in two or four
```

Out[38]:

|            | A         | B         | C         | D         | E     |
|------------|-----------|-----------|-----------|-----------|-------|
| 2013-01-01 | -0.564503 | -1.195000 | 0.441140  | -0.291384 | one   |
| 2013-01-02 | -0.893038 | -0.372430 | -0.436048 | -0.545141 | one   |
| 2013-01-04 | 0.447390  | 0.205224  | -0.614702 | 0.212125  | three |
| 2013-01-06 | 1.000494  | -0.462591 | -0.638053 | 0.481087  | three |

Filtering using query() method

In [52]:

```
import numpy as np
import pandas as pd

dates = pd.date_range('20130101', periods=6)
df = pd.DataFrame(np.random.randn(6,4), index=dates, columns=list('ABCD'))
df
```

Out[52]:

|            | A         | B         | C         | D         |
|------------|-----------|-----------|-----------|-----------|
| 2013-01-01 | -0.683630 | -0.961407 | 1.052018  | 0.512382  |
| 2013-01-02 | -0.755844 | 1.672273  | 0.365384  | 1.030149  |
| 2013-01-03 | -0.532466 | 0.616725  | 1.066731  | -1.056756 |
| 2013-01-04 | 0.153878  | -1.270198 | 0.346587  | 0.150570  |
| 2013-01-05 | 1.645820  | 0.817894  | 0.168950  | -0.833291 |
| 2013-01-06 | -0.768034 | 1.145290  | -0.612168 | 0.325120  |

In [58]:

```
df.query('A > C')
```

Out[58]:

|            | A       | B        | C       | D         |
|------------|---------|----------|---------|-----------|
| 2013-01-05 | 1.64582 | 0.817894 | 0.16895 | -0.833291 |

In [69]:

```
df.query('A > 0')
```

Out[69]:

|            | A        | B         | C        | D         |
|------------|----------|-----------|----------|-----------|
| 2013-01-04 | 0.153878 | -1.270198 | 0.346587 | 0.150570  |
| 2013-01-05 | 1.645820 | 0.817894  | 0.168950 | -0.833291 |

In [68]:

```
df.query('A > 0 & A < 1')
```

Out[68]:

|            | A        | B         | C        | D       |
|------------|----------|-----------|----------|---------|
| 2013-01-04 | 0.153878 | -1.270198 | 0.346587 | 0.15057 |

In [65]:

```
df.query('A > B | A > C') # where A is greater than B or A is greater than C
```

Out[65]:

|            | A         | B         | C        | D         |
|------------|-----------|-----------|----------|-----------|
| 2013-01-01 | -0.683630 | -0.961407 | 1.052018 | 0.512382  |
| 2013-01-04 | 0.153878  | -1.270198 | 0.346587 | 0.150570  |
| 2013-01-05 | 1.645820  | 0.817894  | 0.168950 | -0.833291 |

## Obtaining columns with partial column labels

[\[back to top\]](#)

In [40]:

```
import pandas as pd

df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                   'key2' : ['one', 'two', 'one', 'two', 'one'],
                   'data1' : np.random.randn(5),
                   'data2' : np.random.randn(5)})

df
```

Out[40]:

|   | data1     | data2     | key1 | key2 |
|---|-----------|-----------|------|------|
| 0 | 0.740132  | 1.220225  | a    | one  |
| 1 | 1.223369  | 0.830145  | a    | two  |
| 2 | -2.046775 | -1.149754 | b    | one  |
| 3 | -3.068403 | 0.675124  | b    | two  |
| 4 | -0.467039 | 0.640412  | a    | one  |

In [42]:

```
df.filter(like='data')
```

Out[42]:

|   | data1     | data2     |
|---|-----------|-----------|
| 0 | 0.740132  | 1.220225  |
| 1 | 1.223369  | 0.830145  |
| 2 | -2.046775 | -1.149754 |
| 3 | -3.068403 | 0.675124  |
| 4 | -0.467039 | 0.640412  |

## Getting Value Counts

It is so tempting to use the `groupby()` function or `pivot_table`, but most of the time, `value_counts()` function is all we need.

[\[back to top\]](#)

In [1]:

```
import pandas as pd
data = pd.DataFrame({'group': ['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'c'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[1]:

|   | group | ounces |
|---|-------|--------|
| 0 | a     | 4.0    |
| 1 | a     | 3.0    |
| 2 | a     | 12.0   |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |

In [3]:

```
data.group.value_counts()
```

Out[3]:

```
a    3
b    3
c    3
Name: group, dtype: int64
```

## Getting Cumulative Sum

[\[back to top\]](#)

In [3]:

```
data.group.value_counts().cumsum()
```

Out[3]:

```
a    3
b    6
c    9
dtype: int64
```

**Pivot table example - see also this excellent [article \(http://pbpython.com/pandas-pivot-table-explained.html\)](http://pbpython.com/pandas-pivot-table-explained.html) on pivot tables**

[\[back to top\]](#)

In [51]:

```
import pandas as pd
data = pd.DataFrame({'group': ['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'c'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[51]:

|   | group | ounces |
|---|-------|--------|
| 0 | a     | 4.0    |
| 1 | a     | 3.0    |
| 2 | a     | 12.0   |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |

### Calculating the means of each group

In [3]:

```
data.pivot_table(values='ounces',index='group',aggfunc=np.mean)
```

Out[3]:

```
group
a      6.333333
b      7.166667
c      4.666667
Name: ounces, dtype: float64
```

### Getting counts by group

In [5]:

```
data.pivot_table(values='ounces',index='group',aggfunc='count')
```

Out[5]:

```
group
a      3
b      3
c      3
Name: ounces, dtype: int64
```

### Getting cumulative sum or running total of the group counts

In [6]:

```
data.pivot_table(values='ounces',index='group',aggfunc='count').cumsum()
```

Out[6]:

```
group
a      3
b      6
c      9
Name: ounces, dtype: int64
```

In [14]:

```
import pandas as pd

# sample data can be downloaded here: http://pbpython.com/extras/sales-funnel.xlsx
df = pd.read_csv('/home/pybokeh/Downloads/sales-funnel.csv')
```



In [5]:

df.head()

Out[5]:

|   | Account | Name                         | Rep           | Manager      | Product     | Quantity | Price | Status    |
|---|---------|------------------------------|---------------|--------------|-------------|----------|-------|-----------|
| 0 | 714466  | Trantow-Barrows              | Craig Booker  | Debra Henley | CPU         | 1        | 30000 | presented |
| 1 | 714466  | Trantow-Barrows              | Craig Booker  | Debra Henley | Software    | 1        | 10000 | presented |
| 2 | 714466  | Trantow-Barrows              | Craig Booker  | Debra Henley | Maintenance | 2        | 5000  | pending   |
| 3 | 737550  | Fritsch, Russel and Anderson | Craig Booker  | Debra Henley | CPU         | 1        | 35000 | declined  |
| 4 | 146832  | Kiehn-Spinka                 | Daniel Hilton | Debra Henley | CPU         | 2        | 65000 | won       |

In [15]:

```
by_rep = pd.pivot_table(data=df, index='Rep', columns='Product', values='Quantity', aggfunc='count', fill_value=0)
by_rep
```

Out[15]:

| Product       | CPU | Maintenance | Monitor | Software |
|---------------|-----|-------------|---------|----------|
| Rep           |     |             |         |          |
| Cedric Moss   | 2   | 1           | 0       | 1        |
| Craig Booker  | 2   | 1           | 0       | 1        |
| Daniel Hilton | 2   | 0           | 0       | 1        |
| John Smith    | 1   | 1           | 0       | 0        |
| Wendy Yule    | 2   | 1           | 1       | 0        |

## Creating a new column based on a grouping using transform() method

[\[back to top\]](#)

Let's get sample data from dplython library which is a clone of dplyr

In [6]:

```
from dplython import *
```

In [7]:

```
diamonds.head()
```

Out[7]:

|   | Unnamed: 0 | carat | cut     | color | clarity | depth | table | price | x    | y    | z    |
|---|------------|-------|---------|-------|---------|-------|-------|-------|------|------|------|
| 0 | 1          | 0.23  | Ideal   | E     | SI2     | 61.5  | 55.0  | 326   | 3.95 | 3.98 | 2.43 |
| 1 | 2          | 0.21  | Premium | E     | SI1     | 59.8  | 61.0  | 326   | 3.89 | 3.84 | 2.31 |
| 2 | 3          | 0.23  | Good    | E     | VS1     | 56.9  | 65.0  | 327   | 4.05 | 4.07 | 2.31 |
| 3 | 4          | 0.29  | Premium | I     | VS2     | 62.4  | 58.0  | 334   | 4.20 | 4.23 | 2.63 |
| 4 | 5          | 0.31  | Good    | J     | SI2     | 63.3  | 58.0  | 335   | 4.34 | 4.35 | 2.75 |

**Goal:** Create a column that contains the difference between a diamond's carat and the average of the carats for *that diamond's* color.

Key take-away is to use the transform method of your specific column of the groupby object.

In [21]:

```
(diamonds.assign(carat_diff=diamonds
                  .groupby('color')['carat']
                  .transform(lambda x:x-x.mean())
                  )
  .sort_values(by=['color','carat_diff'])
  .head(10)
)
```

Out[21]:

|       | Unnamed: 0 | carat | cut       | color | clarity | depth | table | price | x    | y    | z    | carat_diff |
|-------|------------|-------|-----------|-------|---------|-------|-------|-------|------|------|------|------------|
| 31597 | 31598      | 0.20  | Ideal     | D     | VS2     | 61.5  | 57.0  | 367   | 3.81 | 3.77 | 2.33 | -0.457795  |
| 31600 | 31601      | 0.20  | Premium   | D     | VS2     | 62.3  | 60.0  | 367   | 3.73 | 3.68 | 2.31 | -0.457795  |
| 31601 | 31602      | 0.20  | Premium   | D     | VS2     | 61.7  | 60.0  | 367   | 3.77 | 3.72 | 2.31 | -0.457795  |
| 38276 | 38277      | 0.21  | Premium   | D     | VS2     | 61.6  | 59.0  | 386   | 3.82 | 3.78 | 2.34 | -0.447795  |
| 38277 | 38278      | 0.21  | Premium   | D     | VS2     | 60.6  | 60.0  | 386   | 3.85 | 3.81 | 2.32 | -0.447795  |
| 38278 | 38279      | 0.21  | Premium   | D     | VS2     | 59.1  | 62.0  | 386   | 3.89 | 3.86 | 2.29 | -0.447795  |
| 38279 | 38280      | 0.21  | Premium   | D     | VS2     | 58.3  | 59.0  | 386   | 3.96 | 3.93 | 2.30 | -0.447795  |
| 54    | 55         | 0.22  | Premium   | D     | VS2     | 59.3  | 62.0  | 404   | 3.91 | 3.88 | 2.31 | -0.437795  |
| 28    | 29         | 0.23  | Very Good | D     | VS2     | 60.5  | 61.0  | 357   | 3.96 | 3.97 | 2.40 | -0.427795  |
| 34    | 35         | 0.23  | Very Good | D     | VS1     | 61.9  | 58.0  | 402   | 3.92 | 3.96 | 2.44 | -0.427795  |

For those curious how to do this using dplyr / dplython:

In [20]:

```
# Using dplyr/dplython syntax
(diamonds >>
  group_by(X.color) >>
  mutate(carat_diff = X.carat - X.carat.mean()) >>
  ungroup() >>
  arrange(X.color, X.carat_diff) >>
  head(10)
)
```

/home/pybokeh/envs/jupyter/lib/python3.5/site-packages/dplython/dplython.py:379: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-vers>  
df[key] = val.applyFcns(df)/home/pybokeh/envs/jupyter/lib/python3.5/site-packages/dplython/dplython.py:429: FutureWarning: sort(columns=....)  
return lambda df: DplyFrame(df.sort(names))

Out[20]:

|       | Unnamed: 0 | carat | cut       | color | clarity | depth | table | price | x    | y    | z    | carat_diff |
|-------|------------|-------|-----------|-------|---------|-------|-------|-------|------|------|------|------------|
| 25981 | 31598      | 0.20  | Ideal     | D     | VS2     | 61.5  | 57.0  | 367   | 3.81 | 3.77 | 2.33 | -0.457795  |
| 25982 | 31601      | 0.20  | Premium   | D     | VS2     | 62.3  | 60.0  | 367   | 3.73 | 3.68 | 2.31 | -0.457795  |
| 25983 | 31602      | 0.20  | Premium   | D     | VS2     | 61.7  | 60.0  | 367   | 3.77 | 3.72 | 2.31 | -0.457795  |
| 26960 | 38277      | 0.21  | Premium   | D     | VS2     | 61.6  | 59.0  | 386   | 3.82 | 3.78 | 2.34 | -0.447795  |
| 26961 | 38278      | 0.21  | Premium   | D     | VS2     | 60.6  | 60.0  | 386   | 3.85 | 3.81 | 2.32 | -0.447795  |
| 26962 | 38279      | 0.21  | Premium   | D     | VS2     | 59.1  | 62.0  | 386   | 3.89 | 3.86 | 2.29 | -0.447795  |
| 26963 | 38280      | 0.21  | Premium   | D     | VS2     | 58.3  | 59.0  | 386   | 3.96 | 3.93 | 2.30 | -0.447795  |
| 22409 | 55         | 0.22  | Premium   | D     | VS2     | 59.3  | 62.0  | 404   | 3.91 | 3.88 | 2.31 | -0.437795  |
| 22404 | 29         | 0.23  | Very Good | D     | VS2     | 60.5  | 61.0  | 357   | 3.96 | 3.97 | 2.40 | -0.427795  |
| 22405 | 35         | 0.23  | Very Good | D     | VS1     | 61.9  | 58.0  | 402   | 3.92 | 3.96 | 2.44 | -0.427795  |

Here's a [comparison \(http://nbviewer.jupyter.org/github/pybokeh/jupyter\\_notebooks/blob/master/dplython/dplython\\_example.ipynb\)](http://nbviewer.jupyter.org/github/pybokeh/jupyter_notebooks/blob/master/dplython/dplython_example.ipynb) I made between dplython and pandas.

Percent of rows

[\[back to top\]](#)

In [24]:

```
perc_of_rows = by_rep.apply(lambda x : x / x.sum() * 100, axis='columns') # or axis=1
perc_of_rows
```

Out[24]:

| Product       | CPU       | Maintenance | Monitor | Software  |
|---------------|-----------|-------------|---------|-----------|
| Rep           |           |             |         |           |
| Cedric Moss   | 50.000000 | 25          | 0       | 25.000000 |
| Craig Booker  | 50.000000 | 25          | 0       | 25.000000 |
| Daniel Hilton | 66.666667 | 0           | 0       | 33.333333 |
| John Smith    | 50.000000 | 50          | 0       | 0.000000  |
| Wendy Yule    | 50.000000 | 25          | 25      | 0.000000  |

### Percent of columns

[\[back to top\]](#)

In [22]:

```
perc_of_columns = by_rep.apply(lambda x : x / x.sum() * 100, axis='index') # or axis=0
perc_of_columns
```

Out[22]:

| Product       | CPU       | Maintenance | Monitor | Software  |
|---------------|-----------|-------------|---------|-----------|
| Rep           |           |             |         |           |
| Cedric Moss   | 22.222222 | 25          | 0       | 33.333333 |
| Craig Booker  | 22.222222 | 25          | 0       | 33.333333 |
| Daniel Hilton | 22.222222 | 0           | 0       | 33.333333 |
| John Smith    | 11.111111 | 25          | 0       | 0.000000  |
| Wendy Yule    | 22.222222 | 25          | 100     | 0.000000  |

## Transpose a data frame

[\[back to top\]](#)

Let's say you have a data frame with several columns and having a hard time scrolling to see what the data looks like. Transpose it!

In [2]:

```
import pandas as pd

df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                  'key2' : ['one', 'two', 'one', 'two', 'one'],
                  'data1' : np.random.randn(5),
                  'data2' : np.random.randn(5)})

df
```

Out[2]:

|   | data1     | data2     | key1 | key2 |
|---|-----------|-----------|------|------|
| 0 | -0.548420 | 0.381045  | a    | one  |
| 1 | -0.327573 | 0.273854  | a    | two  |
| 2 | -1.356345 | -0.045686 | b    | one  |
| 3 | -0.476385 | 1.075987  | b    | two  |
| 4 | -0.923575 | -0.973773 | a    | one  |

In [3]:

```
df.head(1).transpose()
```

Out[3]:

|       | 0          |
|-------|------------|
| data1 | -0.5484204 |
| data2 | 0.3810447  |
| key1  | a          |
| key2  | one        |

You can also do `df.dtypes` to accomplish similar output, but it's nice to see actual data along with the columns.

In [4]:

```
df.dtypes
```

Out[4]:

```
data1    float64
data2    float64
key1      object
key2      object
dtype: object
```

## Converting a data frame index to a column

[\[back to top\]](#)

In [49]:

```
import pandas as pd
data = pd.DataFrame({'group': ['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'c'],
                    'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[49]:

|   | group | ounces |
|---|-------|--------|
| 0 | a     | 4.0    |
| 1 | a     | 3.0    |
| 2 | a     | 12.0   |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |

In [50]:

```
data.reset_index(level=0, inplace=True)
data
```

Out[50]:

|   | index | group | ounces |
|---|-------|-------|--------|
| 0 | 0     | a     | 4.0    |
| 1 | 1     | a     | 3.0    |
| 2 | 2     | a     | 12.0   |
| 3 | 3     | b     | 6.0    |
| 4 | 4     | b     | 7.5    |
| 5 | 5     | b     | 8.0    |
| 6 | 6     | c     | 3.0    |
| 7 | 7     | c     | 5.0    |
| 8 | 8     | c     | 6.0    |

## How to set the data frame's index to be one of the columns

[\[back to top\]](#)

In [6]:

```
import pandas as pd
col = ['a', 'b', 'c']
data = pd.DataFrame([[1,2,3],[10,11,12],[20,21,22]],columns=col)
data
```

Out[6]:

|   | a  | b  | c  |
|---|----|----|----|
| 0 | 1  | 2  | 3  |
| 1 | 10 | 11 | 12 |
| 2 | 20 | 21 | 22 |

In [7]:

```
data = data.set_index('a')
data
```

Out[7]:

|    | b  | c  |
|----|----|----|
| a  |    |    |
| 1  | 2  | 3  |
| 10 | 11 | 12 |
| 20 | 21 | 22 |

To remove the index name ("a"), do:

In [8]:

```
data.index.name = None
data
```

Out[8]:

|    | b  | c  |
|----|----|----|
| 1  | 2  | 3  |
| 10 | 11 | 12 |
| 20 | 21 | 22 |

## How to add or fill in missing dates

[\[back to top\]](#)

In [1]:

```
import pandas as pd

idx = pd.date_range('09-01-2013', '09-30-2013')

s = pd.Series({'09-02-2013': 2,
               '09-03-2013': 10,
               '09-06-2013': 5,
               '09-07-2013': 1})
s
```

Out[1]:

```
09-02-2013    2
09-03-2013   10
09-06-2013    5
09-07-2013    1
dtype: int64
```

From above, we have holes in our data. Let's fill in those missing holes using the idx date\_range we specified above.

In [63]:

```
s.index = pd.DatetimeIndex(s.index)

s = s.reindex(idx, fill_value=0)
s
```

Out[63]:

```
2013-09-01    0
2013-09-02    2
2013-09-03   10
2013-09-04    0
2013-09-05    0
2013-09-06    5
2013-09-07    1
2013-09-08    0
2013-09-09    0
2013-09-10    0
2013-09-11    0
2013-09-12    0
2013-09-13    0
2013-09-14    0
2013-09-15    0
2013-09-16    0
2013-09-17    0
2013-09-18    0
2013-09-19    0
2013-09-20    0
2013-09-21    0
2013-09-22    0
2013-09-23    0
2013-09-24    0
2013-09-25    0
2013-09-26    0
2013-09-27    0
2013-09-28    0
2013-09-29    0
2013-09-30    0
Freq: D, dtype: int64
```

## How to connect and query against a database registered as an ODBC data source (Windows)

[\[back to top\]](#)

In [ ]:

```
import pyodbc          # If using PostgreSQL, MySQL, SQLite, Redshift, MS SQL Server, or Oracle, use db.py instead
import pandas as pd
from getpass import getpass # Module that will create a text input widget AND mask your password

userid = 'your_userid'
pw = getpass(prompt='Enter your password: ')

cnxn_string = 'DSN=your_dsn;UID=' + userid + ';PWD=' + pw

cnxn = pyodbc.connect(cnxn_string)
cursor = cnxn.cursor()

sql = """
SELECT * from your_table...
"""

df = pd.read_sql(sql, cnxn, index_col=None, parse_dates={'some_column':"%Y-%m-%d"})

# Close connections
cursor.close()
cnxn.close()
```

Also check out Yhat's [db.py \(https://github.com/yhat/db.py\)](https://github.com/yhat/db.py). Currently, doesn't support IBM DB2. So I'm stuck with pyodbc at work for now.

## How to convert data in wide format to long format using pd.melt()

[\[back to top\]](#)

When using `ggplot` (<https://github.com/yhat/ggplot>) to plot data where you have multiple data series you want to plot, you need to convert data in wide format to long format.

In [7]:

```
from pandas_datareader import data as web
from datetime import datetime
import pandas as pd

stocks = {'tech': ['GOOGL', 'MSFT', 'LNKD', 'YHOO', 'FB', 'HPQ', 'AMZN'],
          'auto': ['TM', 'F', 'GM', 'HMC', 'NSANY', 'HYMTF'],
          'housing': ['HD', 'WMT', 'LOW']}

start_date = datetime(2014,1,1)
end_date = datetime(2014, 11, 30)

def get_px(stock, start, end):
    return web.get_data_yahoo(stock, start, end)['Adj Close']

df = pd.DataFrame({n: get_px(n, start_date, end_date) for n in stocks['tech']})
```

Below is what data looks like in 'wide' format:

In [8]:

```
df.head(10)
```

Out[8]:

|            | AMZN       | FB        | GOOGL      | HPQ       | LNKD       | MSFT      | YHOO      |
|------------|------------|-----------|------------|-----------|------------|-----------|-----------|
| Date       |            |           |            |           |            |           |           |
| 2014-01-02 | 397.970001 | 54.709999 | 557.117134 | 26.726783 | 207.639999 | 35.448341 | 39.590000 |
| 2014-01-03 | 396.440002 | 54.560001 | 553.053047 | 27.383841 | 207.419998 | 35.209856 | 40.119999 |
| 2014-01-06 | 393.630005 | 57.200001 | 559.219245 | 27.335529 | 203.919998 | 34.465786 | 39.930000 |
| 2014-01-07 | 398.029999 | 57.919998 | 570.000007 | 27.229239 | 209.639999 | 34.732887 | 40.919998 |
| 2014-01-08 | 401.920013 | 58.230000 | 571.186226 | 26.523869 | 209.059998 | 34.112826 | 41.020000 |
| 2014-01-09 | 401.010010 | 57.220001 | 565.685676 | 26.678471 | 215.250000 | 33.893420 | 40.919998 |
| 2014-01-10 | 397.660004 | 57.939999 | 565.655647 | 26.765435 | 218.750000 | 34.379931 | 41.230000 |
| 2014-01-13 | 390.980011 | 55.910000 | 562.052070 | 27.171264 | 213.539993 | 33.368755 | 39.990002 |
| 2014-01-14 | 397.540009 | 57.740002 | 575.275276 | 27.876635 | 216.220001 | 34.131905 | 41.139999 |
| 2014-01-15 | 395.869995 | 57.599998 | 574.884895 | 27.866972 | 216.000000 | 35.066764 | 41.070000 |

In [9]:

```
# Create a column from the index values
df.reset_index(level=0, inplace=True)
df.head(10)
```

Out[9]:

|   | Date       | AMZN       | FB        | GOOGL      | HPQ       | LNKD       | MSFT      | YHOO      |
|---|------------|------------|-----------|------------|-----------|------------|-----------|-----------|
| 0 | 2014-01-02 | 397.970001 | 54.709999 | 557.117134 | 26.726783 | 207.639999 | 35.448341 | 39.590000 |
| 1 | 2014-01-03 | 396.440002 | 54.560001 | 553.053047 | 27.383841 | 207.419998 | 35.209856 | 40.119999 |
| 2 | 2014-01-06 | 393.630005 | 57.200001 | 559.219245 | 27.335529 | 203.919998 | 34.465786 | 39.930000 |
| 3 | 2014-01-07 | 398.029999 | 57.919998 | 570.000007 | 27.229239 | 209.639999 | 34.732887 | 40.919998 |
| 4 | 2014-01-08 | 401.920013 | 58.230000 | 571.186226 | 26.523869 | 209.059998 | 34.112826 | 41.020000 |
| 5 | 2014-01-09 | 401.010010 | 57.220001 | 565.685676 | 26.678471 | 215.250000 | 33.893420 | 40.919998 |
| 6 | 2014-01-10 | 397.660004 | 57.939999 | 565.655647 | 26.765435 | 218.750000 | 34.379931 | 41.230000 |
| 7 | 2014-01-13 | 390.980011 | 55.910000 | 562.052070 | 27.171264 | 213.539993 | 33.368755 | 39.990002 |
| 8 | 2014-01-14 | 397.540009 | 57.740002 | 575.275276 | 27.876635 | 216.220001 | 34.131905 | 41.139999 |
| 9 | 2014-01-15 | 395.869995 | 57.599998 | 574.884895 | 27.866972 | 216.000000 | 35.066764 | 41.070000 |

This is what data looks like in 'long' format:

In [10]:

```
df_long = pd.melt(df, id_vars=['Date']).dropna()
df_long.head()
```

Out[10]:

|   | Date       | variable | value      |
|---|------------|----------|------------|
| 0 | 2014-01-02 | AMZN     | 397.970001 |
| 1 | 2014-01-03 | AMZN     | 396.440002 |
| 2 | 2014-01-06 | AMZN     | 393.630005 |
| 3 | 2014-01-07 | AMZN     | 398.029999 |
| 4 | 2014-01-08 | AMZN     | 401.920013 |

In [11]:

```
df_long.tail()
```

Out[11]:

|      | Date       | variable | value     |
|------|------------|----------|-----------|
| 1605 | 2014-11-21 | YHOO     | 51.040001 |
| 1606 | 2014-11-24 | YHOO     | 51.830002 |
| 1607 | 2014-11-25 | YHOO     | 51.720001 |
| 1608 | 2014-11-26 | YHOO     | 51.930000 |
| 1609 | 2014-11-28 | YHOO     | 51.740002 |

In [12]:

```
df_long.rename(columns={'variable': 'Stock', 'value': 'Price'}, inplace=True)
df_long.head(10)
```

Out[12]:

|   | Date       | Stock | Price      |
|---|------------|-------|------------|
| 0 | 2014-01-02 | AMZN  | 397.970001 |
| 1 | 2014-01-03 | AMZN  | 396.440002 |
| 2 | 2014-01-06 | AMZN  | 393.630005 |
| 3 | 2014-01-07 | AMZN  | 398.029999 |
| 4 | 2014-01-08 | AMZN  | 401.920013 |
| 5 | 2014-01-09 | AMZN  | 401.010010 |
| 6 | 2014-01-10 | AMZN  | 397.660004 |
| 7 | 2014-01-13 | AMZN  | 390.980011 |
| 8 | 2014-01-14 | AMZN  | 397.540009 |
| 9 | 2014-01-15 | AMZN  | 395.869995 |

Now we can plot the stock prices. Since the data is in long format, we can take advantage of ggplot's `color=` parameter, which will magically make line charts with different colors for each stock. If the data was in wide format, you would have had to manually specify all the different stocks to plot or use a loop construct.

[\[back to top\]](#) [\[back to section\]](#)

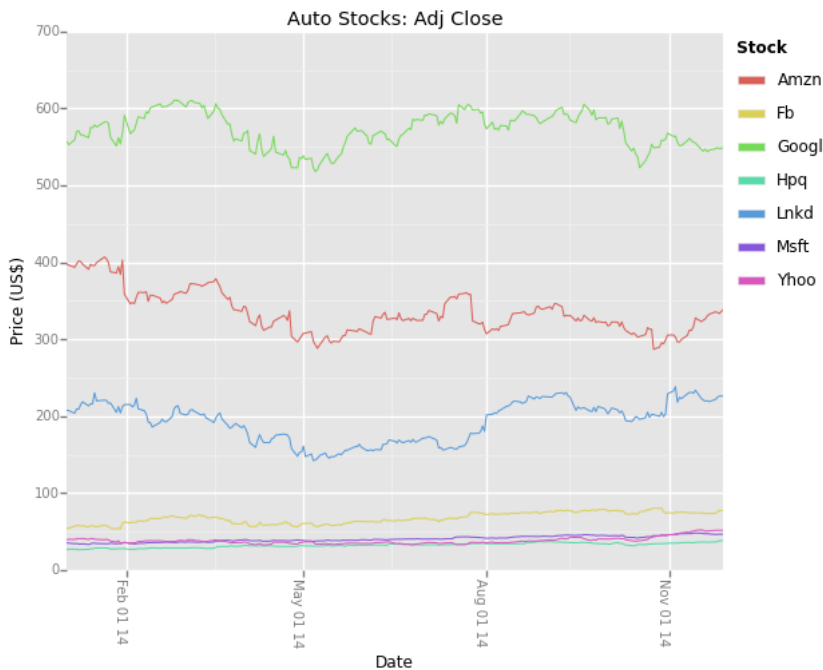


In [13]:

```
%matplotlib inline
import matplotlib.pyplot as plt
from ggplot import *

stock_plot = ggplot(df_long, aes(x='Date', y='Price', color='Stock')) + \
    geom_line() + \
    ylab('Price (US$)') + \
    ggtitle('Auto Stocks: Adj Close') + \
    scale_x_date(labels='%b %d %y',breaks=date_breaks('3 months'))

# Have to use MATPLOTLIB api to rotate x tick labels. Hopefully some day, ggplot will implement a rotation option.
fig = stock_plot.draw()
ax = fig.gca()
labels = ax.get_xticklabels()
for label in labels:
    label.set_rotation(270)
plt.show()
```



**ggplot (<http://blog.yhathq.com/posts/ggplot-for-python.html>) is AWESOME!**

## How to convert data in long format to wide format using pd.pivot()

[\[back to top\]](#)

In [21]:

```
pivoted = df_long.pivot(index='Date', columns='Stock', values='Price')
pivoted.head()
```

Out[21]:

| Stock      | AMZN   | FB    | GOOGL  | HPQ   | LNKD   | MSFT  | YHOO  |
|------------|--------|-------|--------|-------|--------|-------|-------|
| Date       |        |       |        |       |        |       |       |
| 2014-01-02 | 397.97 | 54.71 | 557.12 | 27.17 | 207.64 | 35.91 | 39.59 |
| 2014-01-03 | 396.44 | 54.56 | 553.05 | 27.84 | 207.42 | 35.67 | 40.12 |
| 2014-01-06 | 393.63 | 57.20 | 559.22 | 27.79 | 203.92 | 34.92 | 39.93 |
| 2014-01-07 | 398.03 | 57.92 | 570.00 | 27.68 | 209.64 | 35.19 | 40.92 |
| 2014-01-08 | 401.92 | 58.23 | 571.19 | 26.96 | 209.06 | 34.56 | 41.02 |

## Using category data type to control sort order

[\[back to top\]](#)

In [27]:

```
import pandas as pd
data = pd.DataFrame({'medal': ['bronze', 'silver', 'silver', 'gold', 'bronze', 'bronze', 'gold', 'gold', 'gold'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

Out[27]:

|   | medal  | ounces |
|---|--------|--------|
| 0 | bronze | 4.0    |
| 1 | silver | 3.0    |
| 2 | silver | 12.0   |
| 3 | gold   | 6.0    |
| 4 | bronze | 7.5    |
| 5 | bronze | 8.0    |
| 6 | gold   | 3.0    |
| 7 | gold   | 5.0    |
| 8 | gold   | 6.0    |

In [28]:

```
data["medal"] = data["medal"].astype("category")
data.dtypes
```

Out[28]:

```
medal      category
ounces    float64
dtype: object
```

In [29]:

```
data.sort_values(by=["medal"])
```

Out[29]:

|   | medal  | ounces |
|---|--------|--------|
| 0 | bronze | 4.0    |
| 4 | bronze | 7.5    |
| 5 | bronze | 8.0    |
| 3 | gold   | 6.0    |
| 6 | gold   | 3.0    |
| 7 | gold   | 5.0    |
| 8 | gold   | 6.0    |
| 1 | silver | 3.0    |
| 2 | silver | 12.0   |

**Now define the order of the categorical data using `set_categories()`**

[\[back to top\]](#) [\[back to section\]](#)

In [30]:

```
data["medal"].cat.set_categories(["gold", "silver", "bronze"], inplace=True)
```

In [31]:

```
data.sort_values(by=["medal"])
```

Out[31]:

|   | medal  | ounces |
|---|--------|--------|
| 3 | gold   | 6.0    |
| 6 | gold   | 3.0    |
| 7 | gold   | 5.0    |
| 8 | gold   | 6.0    |
| 1 | silver | 3.0    |
| 2 | silver | 12.0   |
| 0 | bronze | 4.0    |
| 4 | bronze | 7.5    |
| 5 | bronze | 8.0    |

### Merging 2 data frames using merge()

[\[back to top\]](#)

In [72]:

```
import pandas as pd
left = pd.DataFrame({'group': ['a', 'a', 'a', 'b', 'b', 'b', 'c', 'c', 'c'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
left
```

Out[72]:

|   | group | ounces |
|---|-------|--------|
| 0 | a     | 4.0    |
| 1 | a     | 3.0    |
| 2 | a     | 12.0   |
| 3 | b     | 6.0    |
| 4 | b     | 7.5    |
| 5 | b     | 8.0    |
| 6 | c     | 3.0    |
| 7 | c     | 5.0    |
| 8 | c     | 6.0    |

In [76]:

```
right = pd.DataFrame({'label': ['a', 'b', 'c'],
                     'value': ['alpha', 'beta', 'charlie']})
right
```

Out[76]:

|   | label | value   |
|---|-------|---------|
| 0 | a     | alpha   |
| 1 | b     | beta    |
| 2 | c     | charlie |

By default, merge() does an inner join, but you can specify 'left', 'right', or 'outer' joins

In [85]:

```
inner_joined = pd.merge(left, right, how='inner', left_on='group', right_on='label')
inner_joined
```

Out[85]:

|   | group | ounces | label | value   |
|---|-------|--------|-------|---------|
| 0 | a     | 4.0    | a     | alpha   |
| 1 | a     | 3.0    | a     | alpha   |
| 2 | a     | 12.0   | a     | alpha   |
| 3 | b     | 6.0    | b     | beta    |
| 4 | b     | 7.5    | b     | beta    |
| 5 | b     | 8.0    | b     | beta    |
| 6 | c     | 3.0    | c     | charlie |
| 7 | c     | 5.0    | c     | charlie |
| 8 | c     | 6.0    | c     | charlie |

**NOTE:** To merge 2 Series, you have to use `pd.concat()` (<http://stackoverflow.com/questions/18062135/combining-two-series-into-a-dataframe-in-pandas>) instead

For more on merge, join, and concatenate, see the official [docs](http://pandas.pydata.org/pandas-docs/stable/merging.html) (<http://pandas.pydata.org/pandas-docs/stable/merging.html>)

## Finding rows of a data frame containing missing data

[\[back to top\]](#)

In [1]:

```
import numpy as np
import pandas as pd
df = pd.DataFrame([range(3), [0, np.NaN, 0], [0, 0, np.NaN], range(3), range(3)])
df
```

Out[1]:

|   | 0 | 1   | 2   |
|---|---|-----|-----|
| 0 | 0 | 1   | 2   |
| 1 | 0 | NaN | 0   |
| 2 | 0 | 0   | NaN |
| 3 | 0 | 1   | 2   |
| 4 | 0 | 1   | 2   |

In [2]:

```
df.isnull().any(axis=1)
```

Out[2]:

```
0    False
1     True
2     True
3    False
4    False
dtype: bool
```

In [3]:

```
df[df.isnull().any(axis=1)]
```

Out[3]:

|   | 0 | 1   | 2   |
|---|---|-----|-----|
| 1 | 0 | NaN | 0   |
| 2 | 0 | 0   | NaN |

## Converting a data type of a column in a data frame

[\[back to top\]](#)

In [10]:

```
import pandas as pd

a = [['a', '1.2', '4.2'], ['b', '70', '0.03'], ['x', '5', '0']]
df = pd.DataFrame(a, columns=['one', 'two', 'three'])
df
```

Out[10]:

|   | one | two | three |
|---|-----|-----|-------|
| 0 | a   | 1.2 | 4.2   |
| 1 | b   | 70  | 0.03  |
| 2 | x   | 5   | 0     |

In [13]:

df.dtypes

Out[13]:

```
one      object
two      object
three    object
dtype: object
```

In [14]:

```
df[['two', 'three']] = df[['two', 'three']].astype(float)
```

In [15]:

df.dtypes

Out[15]:

```
one      object
two    float64
three    float64
dtype: object
```

## Plotting data frames using MATPLOTLIB version 1.5 and up

[\[back to top\]](#)

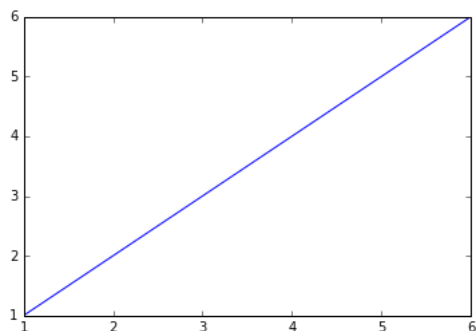
Starting with version 1.5, you can now pass a "data" parameter inside MATPLOTLIB's plot() function

Prior to version 1.5:

In [3]:

```
%matplotlib inline
import pandas as pd
import matplotlib.pyplot as plt

df = pd.DataFrame({"var1": [1, 2, 3, 4, 5, 6], "var2": [1, 2, 3, 4, 5, 6]})
plt.plot(df["var1"], df["var2"])
plt.show()
```



With version 1.5+